

Parallel Concepts

Concurrency

- When multiple operations are making progress within the same time period.
- Usually on the same core/thread.

Parallelism

- When multiple operations are making progress at the same time.
- Usually requires multiple threads/cores.

Process

An instance of the computer that is being executed. These are its components:

- Executable machine language program.
- Block of memory.
- Descriptor of the OS resources allocated to it.
- Security info.
- Information about the state of the process.

Threading

- Threads are contained within processes.
- They allow programmers to divide their programs into independent tasks.
- A stream of instructions that can be scheduled to run independently from its main program.
- The hope is that when one thread blocks because it is waiting for resources, the other can run.

Processes vs Threads

- Threads exist within a process; they're like the children of the process.
- A process has at least one thread.
- If a process has more than one thread, it is multithreaded.
- Starting a thread within a process is known as **forking**.
- Terminating a thread is known as **joining**.
- Both threads and processes are units of execution or **tasks**.
- **Processes do not share memory** (each gets its own block of memory from the system).
- **Threads within a process share memory** (since they are children of the process, they have access to its resources).
- Data stored in a process's memory can be **shared or private**:
 - If **private**, only the thread that owns it can use it.

How is Data Shared?

Shared Memory

- Allows processors to have access to a global address space.

- Multiple processes can operate independently but share the same memory resources.
- Changes in a memory location affected by one task are visible to others.

Uniform Memory Access (UMA)

- The time to access all memory locations is the same for all cores.

Non-Uniform Memory Access (NUMA)

- A memory location a core is directly connected to can be accessed faster than a memory location that must be accessed through another chip.

Task Scheduling

- A **scheduler** is a program that uses a **scheduling policy** to decide which process should run next.
- It uses a **selection function** to make the decision.
- The selection function considers:
 - Resources the process requires.
 - Time the process has been waiting.
 - The process's priority.
- Scheduling policy should try to optimize:
 - **Responsiveness** of interactive processes.
 - **Turnaround time** (the time the user waits for the process to finish).
 - **Resource utilization**.
 - **Fairness** (ensuring each process gets a chance to run).

Some Scheduling Policies

Non-Preemptive Policies

- Each task runs to completion before the next one can run.
- **First In First Out (FIFO)**.
- **Shortest-Job-First (SJF)**.

Preemptive Policies

- **Round-Robin**: Each task is assigned a fixed time before it is required to give way to the next task and move back to the queue.
- **Earliest-Deadline-First**: The process with the closest deadline is picked next.
- **Shortest Remaining Time First**: The process with the shortest remaining time is picked first.

Key Terms

- **Shared resource**: A resource available to all processes in the concurrent program.
- **Critical section**: Sections of code within a process that require access to shared resources. Cannot be executed while another process is in a corresponding section of code.
- **Mutual exclusion**: Requirement that when one process is in a critical section accessing a shared resource, no other process may be in a critical section accessing any of those shared resources.

- **Condition synchronization:** A mechanism ensuring that a process does not proceed until a certain condition is satisfied.
- **Deadlock:** A situation where two or more processes are unable to proceed because each is waiting for another process to act.
- **Livelock:** A situation where two or more processes continuously change their state in response to changes in other processes without making any progress.
- **Race Condition:** A situation where multiple tasks read/write a shared data item, and the result depends on the relative timing of their execution.
- **Starvation:** A situation where a runnable process is overlooked indefinitely by the scheduler.

Dead or Alive(lock)

Concurrent programs must satisfy two properties:

1. **Safety:** The program doesn't enter a bad state.
2. **Liveness:** The program must progress.

Two problems that can occur:

- **Deadlock:** A process is waiting for a shared resource that will never be available (e.g., another process is waiting for this process to act).
- **Livelock:** Multiple processes continuously change state in response to each other without making progress.

Conditions for Deadlock

For deadlock to occur, **four conditions must hold:**

1. **Mutual Exclusion:** The program involves a shared resource protected by mutual exclusion.
2. **Hold While Waiting:** A process can hold a resource while waiting for others.
3. **No Preemption:** The OS cannot force a process to deallocate a resource it holds.
4. **Circular Wait:** P1 is waiting for a resource held by P2, and P2 is waiting for a resource held by P1.

Preventing Deadlock

To prevent deadlock, **prevent at least one of the four conditions** from occurring.

POSIX Threads

A **POSIX thread** is a thread associated with a process's shared resources. Each thread has its own:

- **Stack**
- **Program counter**
- **Registers**
- **Thread ID**

Races

A **race condition** occurs when the **parent process exits before its child threads complete**. This does not allow enough time for child threads to finish execution.

Fixes for Race Conditions

- Best fix for race conditions, use mutual exclusions and join the threads

```
pthread_mutex_t lock;

void* say_something(void *ptr) {
    pthread_mutex_lock(&lock); //this now becomes critical section! it
    uses mutual exclusion
    printf("%s ", (char*)ptr);
    pthread_mutex_unlock(&lock); //end the critical condition
    pthread_exit(0);
}

int main() {
    pthread_t thread_1, thread_2;
    char *msg1 = "Hello ";
    char *msg2 = "World!";
    // create the lock -> error checking?
    pthread_mutex_init(&lock, NULL);
    pthread_create( &thread_1, NULL, say_something, msg1);
    pthread_create( &thread_2, NULL, say_something, msg2);

    // the main thread has to wait for the other threads to terminate
    before it can terminate
    pthread_join(thread_1, NULL);
    pthread_join(thread_2, NULL);

    printf("Done!");
    fflush(stdout);
    pthread_mutex_destroy(&lock);
    exit(0);
}
```

This is conditional synchronization

```
void* say_something(void *ptr) {
    pthread_mutex_lock(&lock); //this now becomes critical section!
    //check on some condition - if it is hello, wait for world....

    if (strcmp("World!", (char*)ptr) == 0) {
        printf("Waiting on condition variable cond1\n");
        if (done == 0) //only wait in the event that you need to...
            pthread_cond_wait(&cond1, &lock);
    } else {
        printf("Signaling condition variable cond1\n");
    }
}
```

```
        done == 1;
        pthread_cond_signal(&cond1);
    }

    printf("%s ", (char*)ptr);
    pthread_mutex_unlock(&lock);
    pthread_exit(0);
}
```

5 - Intro to OpenMP

OpenMP = open Multi-Processing

An api for multithreaded shared parallel programming

OpenMP is:

- higher level than Pthreads
- programmer only states that a block of code is to be executed in parallel
- requires compiler support

Task Parallelisms

- Share the tasks among each core ie on core does the tasks on all data

Data parallelism

- Share the data among each core

OpenMP API

OpenMP is based on directives

OpenMP API components

- Compiler directive
- Runtime library routines
- Environment variables

Fork Join

OpenMP uses the fork join model. The enforces synchronization so every thread must wait till everyone is finished

before proceeding to the next region A group of threads executing the parallel block is known as a team, the original thread is called the master, the children are called slaves

Task parallelism

```
#pragma omp parallel num_threads(4)
{
    int id = omp_get_thread_num();
    printf("T%d:A\n", id);
    printf("T%d:B\n", id);

    if (id == 0)
        printf("T0:special task\n");

    if (id == 1)
        printf("T1:special task\n");

    if(id == 2)
        printf("T2:special task\n");
}

printf("End");
```

Data Parallelism

```
#pragma omp parallel num_threads(2)
{
    int id = get_thread_num()
    int my_a = id * 3;  \\ where you want the thread to start doing work
    int my_b = id * 3 + 3; \\ where it should stop doing work

    printf("T%d will process indexes %d to ");

    for (int index = my_a; index < my_b; index++)
        printf("do work\n");
}

printf("done\n");

return 0;
```

6 - OpenMP Mutexes, Exclusions, and Synchronization

Race Conditions

A **race condition** occurs when multiple threads **simultaneously access and modify shared data**, leading to **unpredictable behavior**.

Example:

```
#pragma omp parallel
{
    global_sum += my_sum; // Potential race condition
}
```

To prevent this, we use **mutual exclusion** techniques.

Barriers

Barriers ensure that all threads reach a synchronization point before continuing execution.

Types of Barriers:

1. **Implicit Barriers** - Automatically added at the end of parallel regions.
2. **Explicit Barriers** - Defined using `#pragma omp barrier`.

Example:

```
#pragma omp parallel
{
    compute_part();
    #pragma omp barrier // Ensures all threads finish before proceeding
    finalize_part();
}
```

Barrier Limitations:

- All threads must encounter the barrier.
 - Conditional execution may lead to **illegal barriers**.
-

nowait Clause

Using `nowait` allows threads to **skip synchronization** when it is unnecessary, improving performance.

Example:

```
#pragma omp single nowait
{
    expensive_task();
}
// Other threads continue execution without waiting.
```

Mutual Exclusion

Mutual exclusion ensures that only **one thread at a time** accesses a critical section.

OpenMP Mutual Exclusion Mechanisms:

1. **Critical Directive** - Ensures exclusive execution.
2. **Atomic Directive** - Ensures atomic updates to a shared variable.
3. **Locks** - Explicit locking mechanisms.

1. Critical Directive

```
#pragma omp critical
{
    shared_var += local_val;
}
```

Named Critical Sections:

```
#pragma omp critical(name1)
x = compute_x();
#pragma omp critical(name2)
y = compute_y();
```

- Allows **simultaneous execution** of **different** critical sections.
-

2. Atomic Directive

`#pragma omp atomic` is **faster** than `critical` for **simple updates**.

```
#pragma omp atomic
sum += value;
```

Supported Operations:

- `x++`, `x--`, `x += expr`, `x = x + expr`
-

3. Locks

Locks **manually enforce** mutual exclusion.

```
#include <omp.h>
static omp_lock_t mylock;

int main() {
    omp_init_lock(&mylock);
```



```
#pragma omp parallel
{
    omp_set_lock(&mylock);
    critical_section();
    omp_unset_lock(&mylock);
}

omp_destroy_lock(&mylock);
return 0;
}
```

Key Lock Functions:

- `omp_init_lock(&lock);`
- `omp_set_lock(&lock);`
- `omp_unset_lock(&lock);`
- `omp_destroy_lock(&lock);`

When to Use Which?

Mechanism	Use Case
Atomic	Single-variable updates (fastest)
Critical	Protects complex code sections
Locks	Fine-grained control over execution

Caveats & Best Practices

1. **Avoid Mixing** different mutual exclusion methods.
2. **Fairness is NOT guaranteed** - Some threads may starve.
3. **Avoid Nesting** critical sections (deadlocks possible).

7 - OpenMP Variable Scope and Reductions

Variable Scope

In OpenMP, variable scope determines which **threads** can access a variable inside a parallel block.

Shared Variables

- Exist in **one memory location**, accessible by all threads.
- Default behavior for variables declared **before** the parallel block.

```
int x = 5;
#pragma omp parallel
{
    // All threads access the same x
}
```

Private Variables

- Each thread gets **its own copy** of the variable.
- Uninitialized unless explicitly set.

```
int y = 5;
#pragma omp parallel private(y)
{
    // Each thread gets its own y (uninitialized)
}
```

Firstprivate Variables

- Like **private**, but **initialized** with the original value.

```
int z = 5;
#pragma omp parallel firstprivate(z)
{
    // Each thread gets its own z, initialized to 5
}
```

Default Clause

Sets the default scope for all variables.

```
int x = 0, y = 0;
#pragma omp parallel num_threads(4) default(none) private(x) shared(y)
{
    x = omp_get_thread_num();
    #pragma omp atomic
    y += x;
}
```

Reductions

Reduction operations allow threads to **aggregate results** safely without manual synchronization.

Syntax

```
#pragma omp parallel reduction(<operator> : <variable list>)
```

Example 1: Summing Across Threads

```
int sum = 0;
#pragma omp parallel reduction(+:sum)
{
    sum += omp_get_thread_num();
}
printf("Total sum = %d", sum);
```

Example 2: Multiple Variables

```
int x = 10, y = 10;
#pragma omp parallel reduction(+:x, y)
{
    x = omp_get_thread_num();
    y = 5;
}
printf("Shared: x=%d, y=%d\n", x, y);
```

Reduction Operations

Operator	Description
+	Summation
*	Multiplication
&	Bitwise AND
	Bitwise OR
^	Bitwise XOR
&&	Logical AND
	Logical OR

Parallel Summation with Reduction

Instead of using a **critical section**, reductions optimize aggregation.



```
double global_sum = 0;
#pragma omp parallel num_threads(4) reduction(+:global_sum)
{
    global_sum += compute_value(omp_get_thread_num());
}
```

Area Under a Curve (Trapezoidal Rule)

Using **reduction** to integrate a function:

```
double global_result = 0.0;
#pragma omp parallel num_threads(4) reduction(+:global_result)
{
    global_result += Local_trap(a, b, n);
}
printf("Approximate area: %f\n", global_result);
```

8 - Work Sharing (Parallel For, Single)

1. Work-Sharing Constructs

- Used to distribute work among threads inside a parallel region.
- **Types:**
 - **for** – Divides loop iterations across threads.
 - **single** – Assigns work to a single thread.
 - **sections** – Splits tasks into sections executed by different threads.
- There is an **implied barrier** at the exit unless **nowait** is specified.

2. Parallel For

- Loop iterations are divided across threads dynamically.
- The loop variable is **private** by default.
- The execution order is **non-deterministic**.

Syntax Options:

1. Inside an existing parallel region:

```
#pragma omp for
for(i = start; i < end; i += step) {
    // Loop body
}
```

2. Creating a parallel region just for the loop:

```
#pragma omp parallel for
for(i = start; i < end; i += step) {
    // Loop body
}
```

Example Without OpenMP Parallelization

```
#pragma omp parallel num_threads(4)
{
    int i, n = omp_get_thread_num();
    for(i=0; i<4; i++)
        printf("T%d: i=%d\n", n , i);
}
```

Each thread executes the whole loop, leading to redundant iterations.

Example With OpenMP Parallel For

```
#pragma omp parallel
{
    int i, n;
    #pragma omp for
    for (i = 0; i < 4; i++) {
        n = omp_get_thread_num();
        printf("T%d: i=%d\n", n, i);
    }
}
```

Iterations are divided among the threads, reducing redundancy.

3. Data Dependency & Loop-Carried Dependencies

- Parallel loops should avoid **loop-carried dependencies** (when one iteration depends on results from another).
- **Example of incorrect parallelization:**

```
fibo[0] = fibo[1] = 1;
#pragma omp parallel for
for (i = 2; i < n; i++)
    fibo[i] = fibo[i-1] + fibo[i-2];
```

This will produce incorrect results because `fibonacci[i-1]` and `fibonacci[i-2]` might not be computed yet.

4. Reduction in Parallel Loops

- Reduction avoids data races when accumulating results.
- **Example: Summing values in an array**

```
double sum = 0.0;
#pragma omp parallel for reduction(+:sum)
for (i = 0; i < n; i++)
    sum += array[i];
```

5. Assigning Work to a Single Thread

- Use `#pragma omp single` for operations that should only be done once.
- **Example:**

```
#pragma omp parallel
{
    printf("Hi from T%d\n", omp_get_thread_num());
    #pragma omp single
    printf("One Hi from T%d\n", omp_get_thread_num());
}
```

Only one thread will execute the `single` block.

9 - Work Sharing (Sections, Scheduling, Ordered Iterations)

1. Parallel Sections

- `#pragma omp sections` allows different sections of code to be executed by different threads.
- **Example:**

```
#pragma omp parallel sections
{
    #pragma omp section
    {
        printf("Section 1 executed by thread %d\n",
            omp_get_thread_num());
    }
    #pragma omp section
```

```

    {
        printf("Section 2 executed by thread %d\n",
omp_get_thread_num());
    }
}

```

- There is an **implicit barrier** at the end of the sections unless `nowait` is used.

2. Loop Scheduling

- The `schedule` clause determines how loop iterations are assigned to threads.

Scheduling Type	Description
<code>static</code>	Equal chunks assigned at compile time.
<code>dynamic</code>	Threads take chunks dynamically.
<code>guided</code>	Starts with large chunks, then reduces.
<code>auto</code>	Compiler decides the best method.

- **Example using dynamic scheduling:**

```

#pragma omp parallel for schedule(dynamic,2)
for(int i = 0; i<8; i++)
    printf("T%d: %d\n", omp_get_thread_num(), i);

```

3. Ordered Iterations

- Ensures that iterations follow a strict order when needed.
- **Example:**

```

#pragma omp for ordered schedule(dynamic)
for(int i=0; i<100; i++) {
    f(a[i]); // Can run in parallel
    #pragma omp ordered
    g(a[i]); // Runs in order
}

```

10 - OpenMP Examples, Functions, SIMD

1. Parallel Matrix Multiplication

```
#pragma omp parallel for collapse(2)
for (i = 0; i < N; i++)
    for (j = 0; j < N; j++) {
        C[i][j] = 0;
        for (k = 0; k < N; k++)
            C[i][j] += A[i][k] * B[k][j];
    }
```

2. Finding the Maximum Value

```
int max_parallel(int *arr){
    int i, m = arr[0];
    #pragma omp parallel for reduction(max:m)
    for (i = 0; i < N; i++)
        if (m < arr[i])
            m = arr[i];
    return m;
}
```

3. Producer-Consumer Model

```
void produce() {
    while (i < NUM_ITEMS) {
        #pragma omp critical(one)
        if (!full) {
            put(item);
            i++;
        }
    }
}

void consume() {
    while (j < NUM_ITEMS) {
        #pragma omp critical(two)
        if (!empty) {
            get();
            j++;
        }
    }
}
```

Ensures only one thread modifies shared data at a time.

Parallel Computing Practice Midterm – Long Answer Solutions

Question 1: Parallelizing Nested Loops


```
### **Given Code:**
```c
for (i = 0; i < N; i++)
 for (j = 0; j < N; j++)
 A[i][j] = max(A[i][j], B[i][j]);
```

### (a) Parallelizing the Code

Using OpenMP, we can parallelize the outer loop to allow multiple threads to work on different rows concurrently.

```
#pragma omp parallel for private(j)
for (i = 0; i < N; i++)
 for (j = 0; j < N; j++)
 A[i][j] = max(A[i][j], B[i][j]);
```

- The `#pragma omp parallel for` ensures each thread handles a different value of `i`.
- `private(j)` ensures each thread has its own copy of `j`.

### (b) Choosing the Best Schedule

- `static` scheduling: Assigns equal chunks of rows to threads. Good if workload is uniform.
- `dynamic` scheduling: Threads request new rows when they finish processing assigned rows. Best for non-uniform workloads.
- `guided` scheduling: Similar to `dynamic`, but chunk sizes decrease over time.

For this case, **static scheduling** is the most efficient since each iteration has equal workload.

```
#pragma omp parallel for schedule(static) private(j)
for (i = 0; i < N; i++)
 for (j = 0; j < N; j++)
 A[i][j] = max(A[i][j], B[i][j]);
```

---

## Question 2: Difference Between Parallel Structures

### Code Snippets & Explanation

#### (a) `#pragma omp master`

```
#pragma omp parallel {
 int n = omp_get_thread_num();
 printf("T%d:A\n", n);
 #pragma omp master
 printf("T%d:X\n", n);
 printf("T%d:B\n", n);
}
```

```

}
printf("Finished");

```

- `#pragma omp master`: Only **one thread (master)** executes `printf("T%d:X\n", n);`.
- All threads execute `printf("T%d:A\n", n);` and `printf("T%d:B\n", n);`.

### (b) `#pragma omp single`

```

#pragma omp parallel {
 int n = omp_get_thread_num();
 printf("T%d:A\n", n);
 #pragma omp single
 printf("T%d:X\n", n);
 printf("T%d:B\n", n);
}
printf("Finished");

```

- `#pragma omp single`: **Only one thread executes** `printf("T%d:X\n", n);`, but it can be any thread, not necessarily the master thread.

### (c) Explicit Check for Thread 0

```

#pragma omp parallel {
 int n = omp_get_thread_num();
 printf("T%d:A\n", n);
 if(omp_get_thread_num() == 0)
 printf("T%d:X\n", n);
 printf("T%d:B\n", n);
}
printf("Finished");

```

- This explicitly checks if the thread number is `0`, similar to `master`, but allows more flexibility.

### Summary of Differences:

- `master`: Only the master thread executes the block.
- `single`: A single (but arbitrary) thread executes the block.
- Explicit check: A thread with a specific ID executes the block.

## Question 3: Parallelizing Loops with Dependencies

### (a) Serial Code:

```

C[0] = 1;
for (i = 1; i < N; i++) {

```

```

 C[i] = C[i - 1];
 for (j = 0; j < N; j++) {
 C[i] *= A[i][j] + B[i][j];
 }
}

```

### Parallelized Version:

- The loop **depends on C[i-1]**, so it **cannot** be fully parallelized.
- However, the inner loop can be parallelized:

```

C[0] = 1;
for (i = 1; i < N; i++) {
 C[i] = C[i - 1];
 #pragma omp parallel for
 for (j = 0; j < N; j++) {
 C[i] *= A[i][j] + B[i][j];
 }
}

```

---

## Question 5: Parallelizing Floyd-Warshall Algorithm

### Given Code:

```

for (k = 0; k < n; k++)
 for (i = 0; i < n; i++)
 for (j = 0; j < n; j++)
 if ((d[i][k] + d[k][j]) < d[i][j])
 d[i][j] = d[i][k] + d[k][j];

```

### Parallelizing It:

Since **d[i][j]** depends on previous iterations of **k**, only the **inner two loops** can be parallelized:

```

for (k = 0; k < n; k++) {
 #pragma omp parallel for collapse(2)
 for (i = 0; i < n; i++)
 for (j = 0; j < n; j++)
 if ((d[i][k] + d[k][j]) < d[i][j])
 d[i][j] = d[i][k] + d[k][j];
}

```

- **collapse(2)**: Merges the two loops so that OpenMP distributes **both i and j** iterations among threads.

---

## Question 6: Explicit OpenMP Parallelization

### Given OpenMP Code:

```
void vector_add(double *a, double *b, double *sum, int n) {
 int i;
 #pragma omp parallel for
 for (i = 0; i < n; i++)
 sum[i] = a[i] + b[i];
}
```

### Manually Managing Threads

Instead of `#pragma omp`, we create threads explicitly:

```
void vector_add(double *a, double *b, double *sum, int n) {
 int TID, TOT;
 #pragma omp parallel private(TID)
 {
 TID = omp_get_thread_num();
 TOT = omp_get_num_threads();
 int range = n / TOT;
 int start = TID * range;
 int end = start + range;

 for (int i = start; i < end; i++) {
 sum[i] = a[i] + b[i];
 }
 }
}
```

- `omp_get_thread_num()`: Each thread gets its unique ID.
- `omp_get_num_threads()`: Gets the total number of threads.
- `range = n / TOT`: Each thread processes an equal chunk.

---

## Code Snippets

---

```
pthread_mutex_t lock;
void* say_something(void *ptr) {
 pthread_mutex_lock(&lock);
 printf("%s ", (char*)ptr);
 pthread_mutex_unlock(&lock);
 pthread_exit(0);
}
```

```

int main() {
 pthread_t t1, t2;
 char *msg1 = "Hello ", *msg2 = "World!";
 pthread_mutex_init(&lock, NULL);
 pthread_create(&t1, NULL, say_something, msg1);
 pthread_create(&t2, NULL, say_something, msg2);
 pthread_join(t1, NULL);
 pthread_join(t2, NULL);
 printf("Done!");
 pthread_mutex_destroy(&lock);
 exit(0);
}

```

### Mutex synchronization for thread safety.

```

#pragma omp parallel num_threads(4)
{
 int id = omp_get_thread_num();
 printf("T%d:A\nT%d:B\n", id, id);
 if (id == 0) printf("T0:special task\n");
 if (id == 1) printf("T1:special task\n");
 if (id == 2) printf("T2:special task\n");
}
printf("End");

```

### Task parallelism in OpenMP.

```

#pragma omp parallel num_threads(2)
{
 int id = omp_get_thread_num();
 int my_a = id * 3, my_b = id * 3 + 3;
 printf("T%d will process indexes %d to %d\n", id, my_a, my_b);
 for (int index = my_a; index < my_b; index++) printf("do work\n");
}
printf("done\n");

```

### Data parallelism using OpenMP.

```

#pragma omp parallel
{
 global_sum += my_sum;
}

```

### Race condition due to unsynchronized access.

```
#pragma omp atomic
sum += value;
```

**Atomic directive ensures safe updates.**

```
#pragma omp critical
{
 shared_var += local_val;
}
```

**Critical section to prevent concurrent access.**

```
#pragma omp parallel for reduction(+:sum)
for (int i = 0; i < n; i++)
 sum += array[i];
```

**Reduction safely aggregates results.**

```
#pragma omp parallel for collapse(2)
for (int i = 0; i < N; i++)
 for (int j = 0; j < N; j++) {
 C[i][j] = 0;
 for (int k = 0; k < N; k++)
 C[i][j] += A[i][k] * B[k][j];
 }
```

**Parallel matrix multiplication using OpenMP.**

```
#pragma omp parallel for schedule(dynamic,2)
for (int i = 0; i < 8; i++)
 printf("T%d: %d\n", omp_get_thread_num(), i);
```

**Dynamic scheduling distributes workload efficiently.**