

COSC 407

Intro to Parallel Computing

Topic 4 – POSIX Threads

COSC 407: Intro to Parallel Computing

1

Outline

Previous pre-recorded lecture (Students' led Q/As):

More on C programming:

- Intro to parallel computing
- Intro to POSIX Threads

Today's topics:

- POSIX Threads – key concepts

Next Lecture:

- Intro to OpenMP

Topic 4: POSIX Threads

COSC 407: Intro to Parallel Computing

2

POSIX Threads

- Key things
 - Parallel concepts with pthreads
 - Thread management
 - Synchronization
 - Mutexes
 - Condition variables

Topic 4: POSIX Threads

COSC 407: Intro to Parallel Computing

3



POSIX Threads

- Threads associated with a process share resources
- Each thread has their own
 - Stack (private variables)
 - Program Counter (PC)
 - Registers
 - Thread ID

Topic 4: POSIX Threads

COSC 407: Intro to Parallel Computing

4



Creating Threads

- To create a thread a `pthread_create` function is used that needs four arguments
 - Thread variable (holds the reference to thread)
 - Thread attribute (specifies the minimum stack size to be used)
 - Function to call when thread starts
 - Arguments to pass to function

```
pthread_t      a_thread;
pthread_attr_t a_thread_attribute; //pthread_attr_default
void thread_function(void *argument);
char          *some_argument;

pthread_create( &a_thread,
               a_thread_attribute,
               (void *)&thread_function,
               (void *)&some_argument
               );
```

Creating Threads

- Threads begin their execution at the function specified in `pthread_create`
- For each thread, we will need to create a thread variable
- Let's consider an example with two threads that will print out a message
 - One thread prints **"Hello "**
 - One thread prints **"World!"**

Example #1

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

void* say_something(void *ptr)
{
    printf("%s ", (char*)ptr);
    return NULL;
}

int main()
{
    pthread_t thread_1, thread_2;

    char *msg1 = "Hello ";
    char *msg2 = "World!";

    pthread_create( &thread_1, NULL, say_something, msg1);
    pthread_create( &thread_2, NULL, say_something, msg2);
    printf("Done!");
    fflush(stdout);
    exit(0);
}
```

- 1 - What does this code do?
- 2 - What is the issue?
 - Run the code multiple times...

Topic 4: POSIX Threads

COSC 407: Intro to Parallel Computing

7

Challenges with Example #1

- Threads execute concurrently
 - There is no guarantee that the first thread reaches the `printf` function prior to the second thread
 - Output could be
 - "World Hello"
 - "Hello World"
 - "World " or " Hello"
 - Nothing.... Why?
- Call to `exit` made by the parent thread in the main block
 - If the parent thread executes the `exit` call prior to either of the child threads executing `printf`, no output will be generated at all.
 - The `exit` function exits the process (releases the task) thus terminating all threads
 - Any thread, parent or child, who calls `exit` can terminate all the other threads along with the process

Topic 4: POSIX Threads

COSC 407: Intro to Parallel Computing

8



Races

- A **race** condition...
 - Exiting before threads exit
 - Time for threads to complete
- We want each child thread to finish before the parent thread
 - Insert a delay in the parent that will give the children time to reach `printf`
 - Ensure that the first child thread reaches `printf` before the second
 - Insert a delay prior to the `pthread_create` call that creates the second thread

Bad Idea... but let's have a look with example 2

Example #2

```
int main()
{
    pthread_t thread_1, thread_2;

    char *msg1 = "Hello ";
    char *msg2 = "World!";

    pthread_create( &thread_1, NULL, say_something, msg1);
    sleep(1);

    pthread_create( &thread_2, NULL, say_something, msg2);
    sleep(1);

    printf("Done!");
    fflush(stdout);
    exit(0);
}
```



Issues with Example #2

- This code doesn't really meet the objectives
 - Not safe to rely on timing delays for synchronization
 - Race condition is still present
 - Sleep function impacts entire process
 - Everything is stalled!
 - Our program just takes longer to run....

11



Allowing Threads to Terminate

- The `pthread_join()` function waits for the thread specified by `thread` to terminate
 - If that thread has already terminated, then `pthread_join()` returns immediately

```
int pthread_join(pthread_t thread, void **retval);
```

https://man7.org/linux/man-pages/man3/pthread_join.3.html

12

Example #3

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

void* say_something(void *ptr)
{
    printf("%s ", (char*)ptr);
    return NULL;
};

int main()
{
    pthread_t thread_1, thread_2;

    char *msg1 = "Hello ";
    char *msg2 = "World!";

    pthread_create( &thread_1, NULL, say_something, msg1);
    pthread_create( &thread_2, NULL, say_something, msg2);

    pthread_join(thread_2, NULL);
    pthread_join(thread_1, NULL);

    printf("Done!");
    fflush(stdout);
    exit(0);
}
```

■ Problems?

- While this happily waits, there is no control over who finishes first!
- How can we sync this so that it works properly??

13



The MUTEX aka Mutual Exclusion

- Threads are lacking Synchronization
 - i.e. who gets to run/access things first?
- Thread synchronization can be achieved using a **Mutex (Mutual Exclusion)**
 - Only one thread at a time can have access to a shared resource
- A Mutex is a lock that is set before using a shared resource and release after using it
 - When the lock is set, no other thread can access the locked region of code (critical code section)
 - Ensures synchronized access of shared resources in the code
 - Can be used to protect access to key resources

14



The MUTEX aka Mutual Exclusion

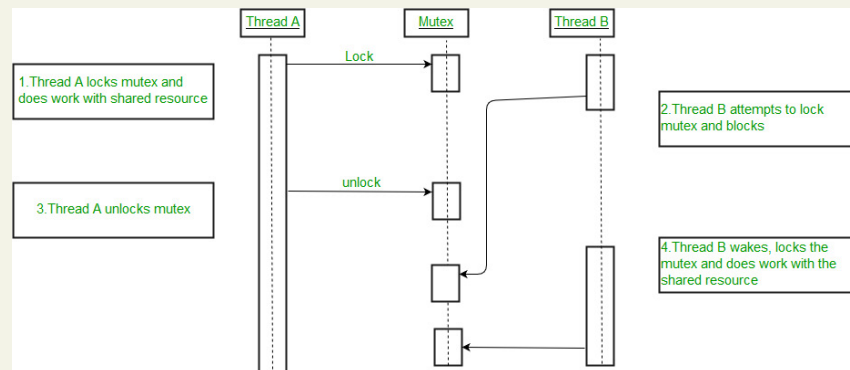
- A mutex is initialized and then a lock is achieved
 - Initializes a mutex
- `pthread_mutex_init(&lock, NULL);`
- Achieve/test lock of the critical code section
- Mutex needs to be released when done

```
void* say_something(void *ptr)
{
    //this now becomes critical section!
    pthread_mutex_lock(&lock);
    printf("%s ", (char*)ptr);
    pthread_mutex_unlock(&lock);
    pthread_exit(0);
}
```

- Mutexes need to be deleted (destroyed) when done with them

15

Mutexes



16

Example #4

```
//get a lock
pthread_mutex_t lock;

void* say_something(void *ptr)
{
    pthread_mutex_lock(&lock); //this now becomes critical section!
    printf("%s ", (char*)ptr);
    pthread_mutex_unlock(&lock);
    pthread_exit(0);
}

int main()
{
    pthread_t thread_1, thread_2;

    char *msg1 = "Hello ";
    char *msg2 = "World!";

    //create the lock -> error checking?
    pthread_mutex_init(&lock, NULL);

    pthread_create(&thread_1, NULL, say_something, msg1);
    pthread_create(&thread_2, NULL, say_something, msg2);

    pthread_join(thread_1, NULL);
    pthread_join(thread_2, NULL);
    printf("Done!");
    fflush(stdout);

    pthread_mutex_destroy(&lock);
    exit(0);
}
```

- Still problems?
 - The thread that gets the lock first, gets to go first

Topic 4: POSIX Threads

COSC 407: Intro to Parallel Computing

17



Conditions

- There are many cases where a thread wishes to check whether a condition is true before continuing its execution
- Not a variable but are used with an associated mutex
- A condition variable is an explicit queue that threads can put themselves on when some state of execution (i.e., some condition) is not as desired (by waiting on the condition);
- When it changes said state, can then wake one (or more) of those waiting threads and thus allow them to continue

`pthread_cond_wait(&cond1, &lock);` If the condition is not true, release lock and wait

`pthread_cond_signal(&cond1);` wake up threads waiting for the condition variable.

Topic 4: POSIX Threads

COSC 407: Intro to Parallel Computing

18

Example 5

```
void* say_something(void *ptr)
{
    pthread_mutex_lock(&lock); //this now becomes critical section!

    //check on some condition - if it is hello, wait for world....
    if (strcmp("World!", (char*)ptr) == 0)
    {
        printf("Waiting on condition variable cond1\n");
        if (done == 0) //only wait in the event that you need to...
            pthread_cond_wait(&cond1, &lock);
    }
    else
    {
        printf("Signaling condition variable cond1\n");
        done = 1;
        pthread_cond_signal(&cond1);
    }
    printf("%s ", (char*)ptr);
    pthread_mutex_unlock(&lock);
    pthread_exit(0);
}
```

Topic 4: POSIX Threads

COSC 407: Intro to Parallel Computing

19

19



Key Functions

- **pthread_create**
 - Create a thread
- **pthread_join**
 - Wait for thread to compete
- **pthread_mutex_init**
 - Create a lock
- **pthread_mutex_lock/pthread_mutex_unlock**
 - Lock and unlock a mutex (if available)
- **pthread_cond_wait**
 - Check on condition
- **pthread_cond_signal**
 - Signal threads waiting on condition

Topic 4: POSIX Threads

COSC 407: Intro to Parallel Computing

20

20

Conclusion/Up Next

- What we covered today (review key concepts):
 - POSIX Threads – key concepts
 - There is a lot of detail here
 - Gives a basic Idea of challenges
 - Will expand on this with OpenMP
- Next Lecture:
 - OpenMP

21

Homework

- Please review
 - POSIX Threads Programming
 - <https://hpc-tutorials.llnl.gov/posix/> (sections 1 – 8)
 - Additional resources on Canvas
 - Have a look at the example code in the course repo (link on Canvas)

22