

COSC 407

Intro to Parallel Computing

MPI: Collective Communication

COSC 407: Intro to Parallel Computing

Outline

Previously:


- Distributed Memory Programming
- Intro to MPI
- MPI: Point to Point communication
 - Example: Area under the curve in MPI
- Dealing with I/O

Today:

- MPI: Collective communication

Previously...

- A **communicator** has n **nodes**, each with **unique rank**.
- The nodes are interconnected by **channels**.
- Each node can **send** and **receive**.
- Two communication scopes: **point-to-point**, or **collective**
- **SPMD** = Single-Program Multiple-Data.
- **Message Matching** means both sender and receiver must execute **matching functions** for sending and receiving.



Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing



Two types of communications

Point-to-point communications:

- Those function which involves two processes (both one to one)
- MPI_Send
 - MPI_Recv
-

Collective Communication:

Those functions that involve all the processes in a communicator

- MPI_Reduce Reduction (all to one)
- MPI_Allreduce Reduction (all to all)
- MPI_Bcast Broadcast (one to all)
- MPI_Scatter Data distribution (one to all)
- MPI_Gather Data concatenation (all to one)
- MPI_Allgather Data concatenation (all to all)

Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

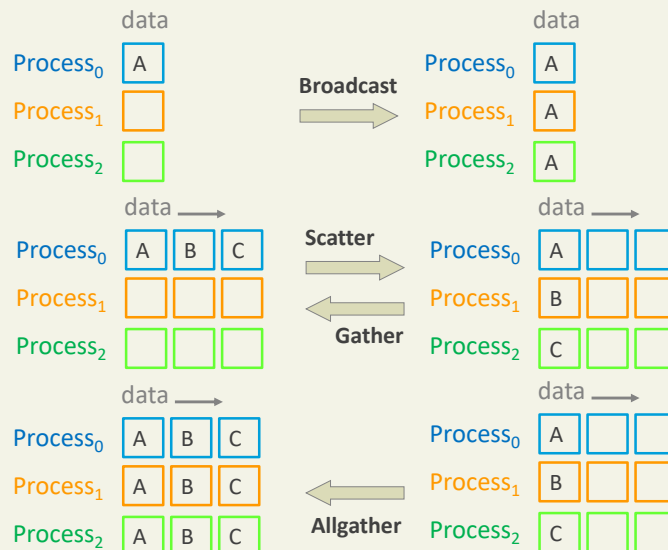
Summary of Collective Comm



Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

Summary of Collective Comm




Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

Reduction: Discussing Trap Rule

- In the previous lecture, we wrote the trapezoid program using point-to-point communication as shown below.
 - Work load is not evenly distributed (process 0 has to do more work)
 - This is ok in case of reading the input
 - But this could be improved for computing the final result.
 - This is point-to-point communication



Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

Remember: Trap Rule: V. 2

```
int main() {
    int my_rank, comm_sz, n, my_n, source;
    double a, b, h, my_a, my_b, my_sum, total_sum;
    //initialize, get rank and comm_sz
    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &comm_sz);
    read_input(my_rank, comm_sz, &a, &b, &n);
    //break down the problem to subproblems
    h = (b-a) / n; // the same for all processes
    my_n = n / comm_sz; // the same for all processes
    my_a = a + my_rank * my_n * h; //unique to each process
    my_b = my_a + my_n * h; //unique to each process
    my_sum = Trap(my_a, my_b, my_n, h); //find partial result from one process
    if (my_rank != 0) { //send my partial result to process 0
        MPI_Send(&my_sum, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD);
    } else { //PROCESS 0 COMBINES THE PARTIAL RESULTS
        total_sum = my_sum;
        for (source = 1; source < comm_sz; source++) {
            MPI_Recv(&my_sum, 1, MPI_DOUBLE, source, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
            total_sum += my_sum;
        }
    }
    if (my_rank == 0) printf("%.15e\n", total_sum);
    MPI_Finalize();
    return 0;
}
```

Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

Tree-structured Communication


- Process 1 sends to 0, 3 sends to 2, ...etc. Processes 0, 2, 4, and 6 add in the received values.
- Processes 2 and 6 send to 0 and 4. Processes 0 and 4 add the received values into their new values.
- Process 4 sends its newest value to process 0. Process 0 adds the received value to its newest value.

The good:

Better workload distribution,
Process 0 is doing less work
(only 3 receives & 3 additions)

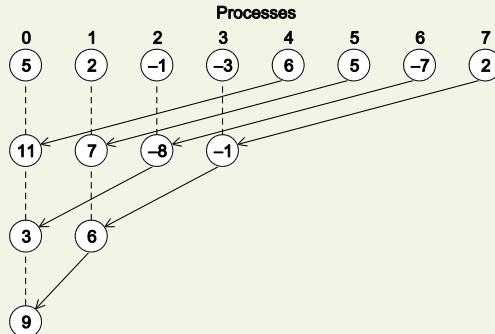
The bad:

- Processes 1,3,5,7 are still not contributing enough
- Requires Extra coding (not easy)!



An Alternative Tree-Structured Global Sum

- Tree structured communication can be planned in different ways. For example, here is another way, and there are many other possibilities.
- Note that this structure suffers from the *same problems as the previous process*.



Is there a way that doesn't involve much coding and testing to see which structure is best?



MPI Reduction

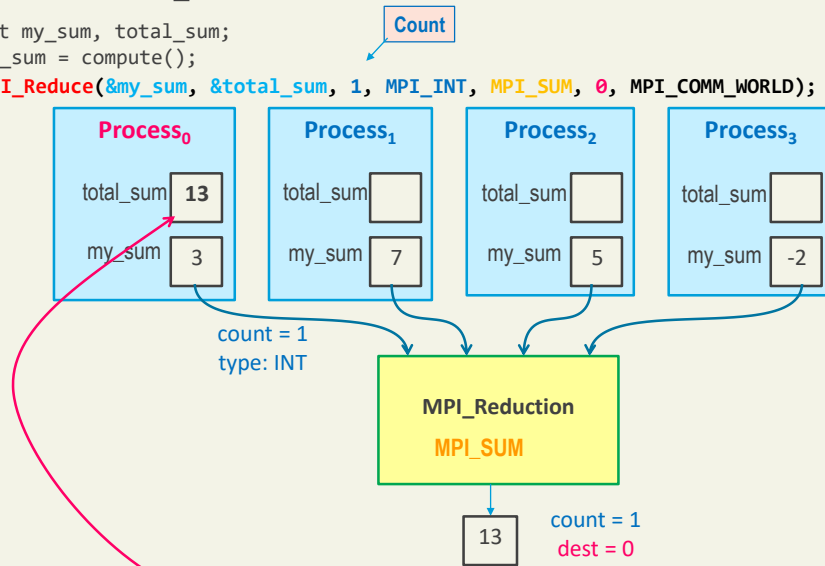
- With the endless possibilities, it's *not reasonable* to ask each MPI programmer to write an optimal global-sum function (much more code and a lot of testing to find the optimal way).
 - Therefore, MPI includes implementations of reduction operators.
 - This means the burden of the extra code is put on the developer of the MPI implementation, rather than the application developer.
 - Two functions that can be used for global reduction:
 - MPI_Reduce Reduction (all to one)
 - MPI_Allreduce Reduction (all to all)
- Obviously, reduction is a form of *collective communication*

MPI_Reduce

```
int MPI_Reduce(  
    void*      input_data_p,      /* in: pointer to input data */  
    void*      output_data_p,     /* out: pointer to output data */  
  
    int        count,             /* in: how many elements in i/p & o/p? */  
    MPI_Datatype datatype,        /* in: datatype of elements in i/p & o/p */  
  
    MPI_Op     operator,          /* in: What type of reduction? */  
  
    int        dest_process,       /* in: ONE process receives results */  
    MPI_Comm   comm               /* in: which communicator? */  
);
```

Example 1

```
int my_sum, total_sum;
my_sum = compute();
MPI_Reduce(&my_sum, &total_sum, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);
```

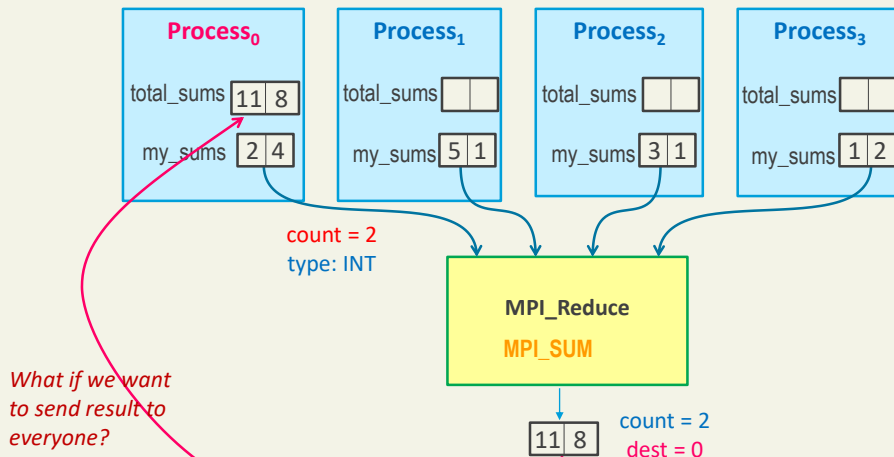


Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

Example 2

```
int my_sums[2], total_sums[2];
my_sums = compute(); //my_sums is an array of 2 elements
MPI_Reduce(&my_sums, &total_sums, 2, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);
```



Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing



MPI_Allreduce

```
int MPI_Allreduce(  
    void*      input_data_p, /* in: address of input data */  
    void*      output_data_p, /* out: address of output data */  
    int        count,         /* in: how many elements in i/p & o/p? */  
    MPI_Datatype datatype,    /* in: datatype of elements in i/p & o/p */  
    MPI_Op      operator,     /* in: What type of reduction? */  
    MPI_Comm    comm         /* in: which communicator? */  
);
```

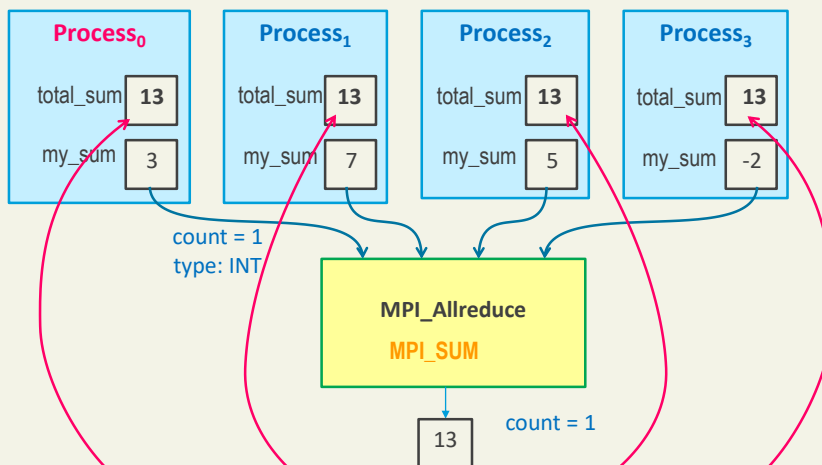
there is no destination process

Same as MPI_Reduce except that result is sent to all processes.

- e.g., all processes need global sum to complete some larger computation.

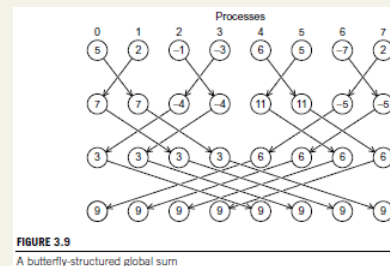
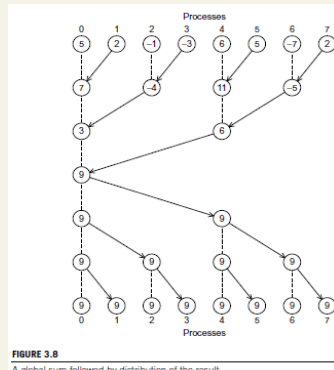
Example 3

```
int my_sum, total_sum; my_sum = compute();  
MPI_Allreduce(&my_sum, &total_sum, 1, MPI_INT, MPI_SUM, MPI_COMM_WORLD);
```



Implementation of MPI_Allreduce

- Note that there are several possible ways of implementing reduction followed by distribution of data to all nodes. Again, it is left to the developer of the MPI implementation to code the best implementation of the operation.



Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

Reduction Operators

- Predefined **reduction operators** in MPI

Operation Value	Meaning
<code>MPI_MAX</code>	Maximum
<code>MPI_MIN</code>	Minimum
<code>MPI_SUM</code>	Sum
<code>MPI_PROD</code>	Product
<code>MPI_LAND</code>	Logical and
<code>MPI_BAND</code>	Bitwise and
<code>MPI_LOR</code>	Logical or
<code>MPI_BOR</code>	Bitwise or
<code>MPI_LXOR</code>	Logical exclusive or
<code>MPI_BXOR</code>	Bitwise exclusive or
<code>MPI_MAXLOC</code>	Maximum and location of maximum
<code>MPI_MINLOC</code>	Minimum and location of minimum

Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

Compare with Version 2

```
int main() {
    //version 2: without reduction
    int my_rank, comm_sz, my_n, source;
    double a, b, h, my_a, my_b, my_sum, total_sum;
    //initialize, get rank and comm_sz
    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &comm_sz);
    read_input(my_rank, comm_sz, &a, &b, &n);
    //break down the problem to subproblems
    h = (b-a) / N; // the same for all processes
    my_n = N / comm_sz; // the same for all processes
    my_a = a + my_rank * my_n * h; //unique to each process
    my_b = my_a + my_n * h; //unique to each process
    my_sum = Trap(my_a, my_b, my_n, h); //find partial result
    //manual reduction (by the programming)
    if (my_rank != 0) { //send my partial result to process 0
        MPI_Send(&my_sum, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD);
    } else { //process 0 combines the partial results
        total_sum = my_sum;
        for (source = 1; source < comm_sz; source++) {
            MPI_Recv(&my_sum, 1, MPI_DOUBLE, source, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
            total_sum += my_sum;
        }
    }
    if (my_rank == 0) printf("%.15e\n", total_sum);
    MPI_Finalize();
    return 0;
}
```

Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

Compare with Version 2

```
int main() {
    //version 3: using reduction
    int my_rank, comm_sz, my_n;
    double a, b, h, my_a, my_b, my_sum, total_sum;
    //initialize, get rank and comm_sz
    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &comm_sz);
    read_input(my_rank, comm_sz, &a, &b, &n);
    //break down the problem to subproblems
    h = (b-a) / N; // the same for all processes
    my_n = n / comm_sz; // the same for all processes
    my_a = a + my_rank * my_n * h; //unique to each process
    my_b = my_a + my_n * h; //unique to each process
    my_sum = Trap(my_a, my_b, my_n, h); //find partial result
    //reduction (by the MPI implementation)
    MPI_Reduce(&my_sum, &total_sum, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);

    if (my_rank == 0) printf("%.15e\n", total_sum);
    MPI_Finalize();
    return 0;
}
```

Ask

Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing



Collective vs. Point-to-Point Communications

1. ALL the processes in the communicator must call the **same collective function**.
 - For example, a program that attempts to match a call to `MPI_Reduce` on one process with a call to `MPI_Recv` on another process is **erroneous**, and, in all likelihood, the program will hang or crash.
2. The arguments passed by each process to an MPI collective communication must be “**compatible**.”
 - For example, if one process passes in 0 as the `dest_process` and another passes in 1, then the outcome of a call to `MPI_Reduce` is **erroneous**, and, once again, the program is likely to hang or crash.
3. In case of `MPI_Reduce`, The `output_data_p` argument is only used on `dest_process`
 - However, all of the processes still need to pass in an actual argument corresponding to `output_data_p`, even if it’s just `NULL`.

Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing



Collective vs. Point-to-Point Communications

4. **BE CAREFUL: Point-to-point communications** are matched on the basis of **tags** and **communicators**.
Collective communications don't use tags. They're matched solely on the basis of the communicator and the **order** in which they're called.
 - Example: Assume each process calls `MPI_Reduce` with `MPI_SUM`, and **destination process** is 0.

Time	Process 0	Process 1	Process 2
0	a = 1; c = 2	a = 1; c = 2	a = 1; c = 2
1	<code>MPI_Reduce(&a, &b, ...)</code>	<code>MPI_Reduce(&c, &d, ...)</code>	<code>MPI_Reduce(&a, &b, ...)</code>
2	<code>MPI_Reduce(&c, &d, ...)</code>	<code>MPI_Reduce(&a, &b, ...)</code>	<code>MPI_Reduce(&c, &d, ...)</code>

- At first glance, it might seem that after the two calls to `MPI_Reduce`, b on Process 0 will be 3 ($b = a + a + a$), and the value of d will be 6 ($d = c + c + c$).
- However, the **names of the memory locations are irrelevant** to the matching of the calls to `MPI_Reduce`. The order of the calls will determine the matching so:
 - b will be $1 + 2 + 1 = 4$ ($b_{\text{Proc0}} = a_{\text{Proc0}} + c_{\text{Proc1}} + a_{\text{Proc2}}$)
 - d will be $2 + 1 + 2 = 5$ ($d_{\text{Proc0}} = c_{\text{Proc0}} + a_{\text{Proc1}} + c_{\text{Proc2}}$)


Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing



Broadcasting

- **Broadcast** is a collective communication in which data belonging to a single process is sent to all of the processes in the communicator.
- The distribution of data could be done in *different ways*. Similarly to reduction, the *MPI implementation decides* the best way of doing this while balancing the workload.
- Conceptually, we are going to think of a broadcast as shown below:



MPI_Bcast

```
int MPI_Bcast(  
    void*      data_p,      /* in/out: address of in/out data */  
  
    int        count,       /* in: how many elements to send/receive */  
    MPI_Datatype datatype,  /* in: datatype of elements */  
  
    int        source_proc, /* in: source process */  
    MPI_Comm   comm        /* in: which communicator? */ );
```

source_proc sends its data
to everyone else

Getting an Input With and Without Broadcast

Without

```
if (my_rank==0){ //only Process 0 reads input then sends it to everyone
    data = readData();
    for each node other than 0:
        MPI_Send(process 0 sends data to node);
} else { //everyone else receive data from process 0
    MPI_Recv(node receives data from 0);
}
```



With

```
if (my_rank==0) //only Process 0 reads input
    data = readData();
MPI_Bcast(process 0 sends data to everyone; everyone receive data from 0)
```

Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

read_input Without Broadcasting

```
void read_input(int my_rank,int comm_sz,double* a_p,double* b_p,int* n_p){
    int dest;
    if (my_rank==0){ //only Process 0 reads input
        printf("Enter a, b, and n\n");
        scanf("%lf %lf %d", a_p, b_p, n_p);
        for (dest = 1; dest < comm_sz; dest++) {
            MPI_Send(a_p, 1, MPI_DOUBLE, dest, 0, MPI_COMM_WORLD);
            MPI_Send(b_p, 1, MPI_DOUBLE, dest, 0, MPI_COMM_WORLD);
            MPI_Send(n_p, 1, MPI_INT, dest, 0, MPI_COMM_WORLD);
        }
    } else { //everyone else receive data from process 0
        MPI_Recv(a_p,1,MPI_DOUBLE,0,0,MPI_COMM_WORLD,MPI_STATUS_IGNORE);
        MPI_Recv(b_p,1,MPI_DOUBLE,0,0,MPI_COMM_WORLD,MPI_STATUS_IGNORE);
        MPI_Recv(n_p,1,MPI_INT, 0,0,MPI_COMM_WORLD,MPI_STATUS_IGNORE);
    }
}
```

Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

read_input WITH Broadcasting

```
void read_input(int my_rank,int comm_sz,double* a_p,double* b_p,int* n_p){
    int dest;
    if (my_rank==0){        //only Process 0 reads input
        printf("Enter a, b, and n\n");
        scanf("%lf %lf %d", a_p, b_p, n_p);
    }
    MPI_Bcast(a_p, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);
    MPI_Bcast(b_p, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);
    MPI_Bcast(n_p, 1, MPI_INT, 0, MPI_COMM_WORLD);
}
```

Everyone MUST execute MPI_Bcast so that
it can be matched.
Process 0 sends, everyone else receive.

Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

Data Partitioning Options

- We can partition the data and send them over to different processes. We might consider these options
 - Block partitioning (THE EASIEST – **We shall use this one**)
 - Assign blocks of consecutive components to each process.
 - Cyclic partitioning
 - Assign components in a round robin fashion.
 - Block-cyclic partitioning
 - Use a cyclic distribution of blocks of components.

Example: partitioning 12-component vector among 3 processes

Process	Components											
	Block				Cyclic				Block-cyclic Blocksize = 2			
0	0	1	2	3	0	3	6	9	0	1	6	7
1	4	5	6	7	1	4	7	10	2	3	8	9
2	8	9	10	11	2	5	8	11	4	5	10	11

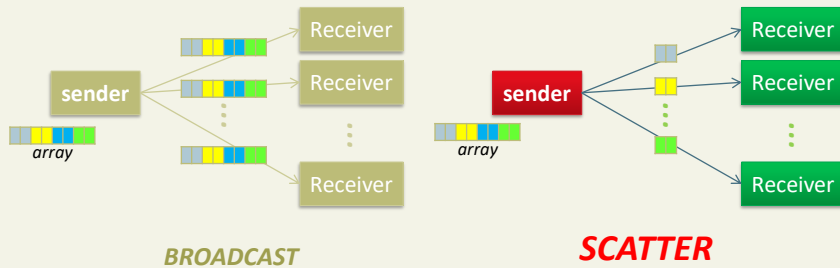
Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing



Scattering

- **MPI Scattering** is a collective routine that similar to broadcasting. The main difference is that:
 - **broadcasting** sends the *same* piece of data to all processes,
 - **scattering** sends *chunks* of an array to different processes.



Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

MPI_Scatter

```

int MPI_Scatter(
    void*          send_buffer_p, /* in: array of data on src process */
    int            send_count,    /* in: # elements to send to EACH process*/
    MPI_Datatype   send_type,    /* in: datatype of elements */

    void*          rcv_buffer_p, /* in: this is where to received data */
    int            rcv_count,    /* in: capacity of rcv buffer*/
    MPI_Datatype   rcv_type,    /* in: datatype of elements */

    int            src_process,  /* in: who sent the data*/
    MPI_Comm       comm         /* in: which communicator? */ );
  
```

MPI_Scatter uses **block partitioning** with the assumption that the length of data is divisible by `comm_sz`. With this assumption,
 $\text{send_count} = \text{send_receive} = N / \text{comm_sz}$

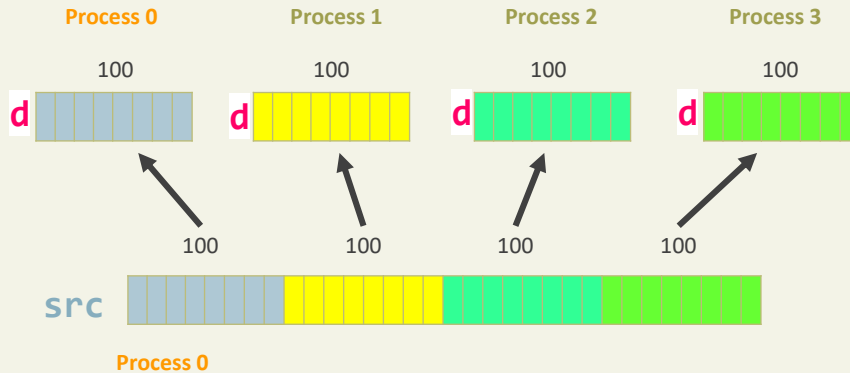
In this course, we will not consider other cases!

Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

Scatter Example

```
MPI_Scatter(src,100,MPI_INT,d,100,MPI_INT,0,MPI_COMM_WORLD)
```



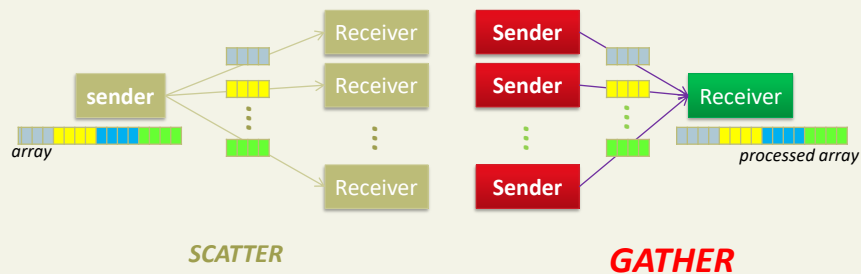
Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing



Gathering

- **MPI Gathering** is a collective routine that performs the inverse of scattering. Gathering combines data chunks from other processes into one big piece of data.
 - The chunks are ordered by the rank of the sending process.



Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

MPI_Gather

```
int MPI_Gather(  
    void*      send_buffer_p, /* in: array of data on src process */  
    int        send_count,    /* in: # elements in send buffer */  
    MPI_Datatype send_type,    /* in: datatype of elements */  
  
    void*      recv_buffer_p, /* in: this is where to received data */  
    int        recv_count,    /* in: #elements received from each process */  
    MPI_Datatype recv_type,    /* in: datatype of elements */  
  
    int        dest_process,   /* in: who collects data */  
    MPI_Comm   comm           /* in: which communicator? */  
);
```

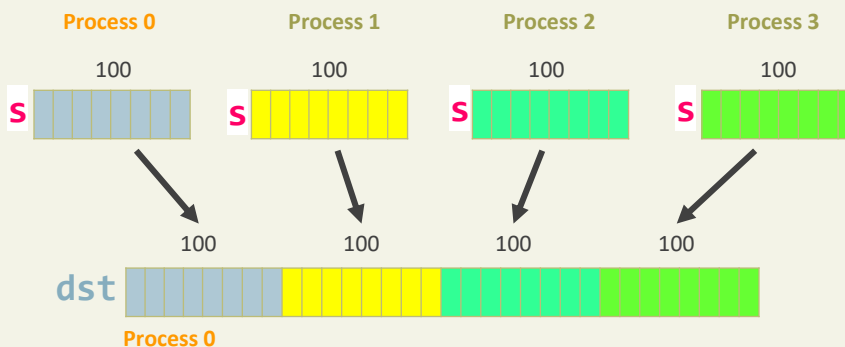
Assuming number of elements in array is N, then
 $\text{send_count} = \text{send_receive} = N / \text{comm_sz}$

Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

Gather Example

```
MPI_Gather(s, 100, MPI_INT, dst, 100, MPI_INT, 0, MPI_COMM_WORLD)
```



Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

MPI_Allgather

Same as gather, but all processes receive a copy of the concatenated vector
– there is no id for receiver process.

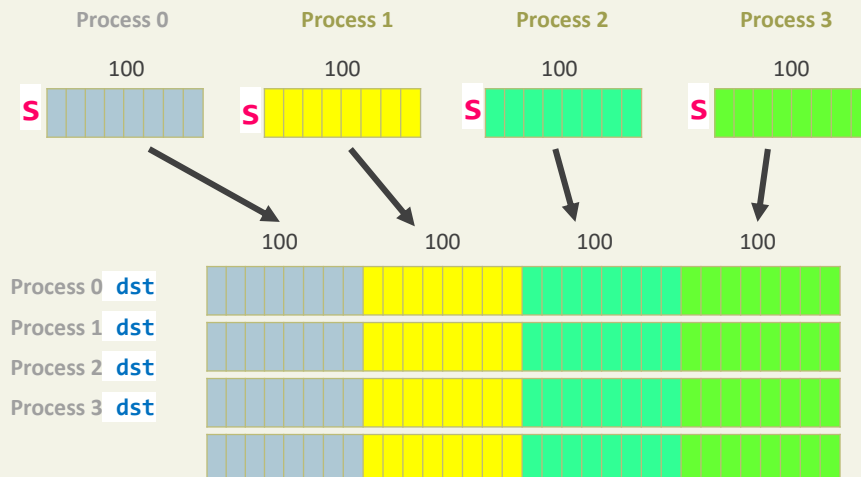
```
int MPI_Allgather(  
    void*      send_buffer_p, /* in: array of data on src process */  
    int        send_count,    /* in: # elements to send to DEST process*/  
    MPI_Datatype send_type,    /* in: datatype of elements */  
  
    void*      recv_buffer_p, /* in: this is where to received data */  
    int        recv_count,    /* in: #elements received from each process*/  
    MPI_Datatype recv_type,    /* in: datatype of elements */  
  
    MPI_Comm    comm          /* in: which communicator? */ );
```

Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

Gather Example

```
MPI_Allgather(s,100,MPI_INT,dst,100,MPI_INT,MPI_COMM_WORLD)
```



Topic 18: MPI Collective Communications

COSC 407: Intro to Parallel Computing

Conclusion

Today:

- MPI: Collective communication

Next day:

- Examples
- MPI Wrap-up
- Course Conclusion